

New algorithms for radio pulsar search

Kendrick M. Smith¹

¹*Perimeter Institute for Theoretical Physics, Waterloo, ON N2L 2Y5, Canada*

(Dated: October 24, 2016)

The computational cost of searching for new pulsars is a limiting factor for upcoming radio telescopes such as SKA. We introduce four new algorithms: an optimal constant-period search, a coherent tree search which permits optimal searching with $\mathcal{O}(1)$ cost per model, a semicoherent search which combines information from coherent subsearches while preserving as much phase information as possible, and a hierarchical search which interpolates between the coherent and semicoherent limits. Taken together, these algorithms improve the computational cost of pulsar search by several orders of magnitude. In this paper, we consider the simple case of a constant-acceleration phase model, but our methods should generalize to more complex search spaces.

I. INTRODUCTION

Pulsar science has been an exceptionally fertile area of astronomy over the last 50 years. Some highlights include first observational evidence for gravitational waves [1], first detection of extrasolar planets [2], and exquisite tests of general relativity from the double pulsar system [3].

Although ≈ 2000 pulsars are known to date, the total observable population is estimated to be larger by a factor ~ 100 [4]. Therefore, finding and timing pulsars is still in a relatively early stage. Upcoming milestones include detection of nanohertz gravitational waves through pulsar timing [5], and discovery of a pulsar-black hole binary system, which would enable new tests of strong gravity [6].

In practice, pulsar searching is very computationally expensive, and tradeoffs between statistical optimality and computational cost are often necessary, particularly for upcoming instruments such as SKA with large numbers of formed beams [7]. Because of this tradeoff, improving computational cost is more than a matter of convenience: it means that searches can be more optimal, thus finding more pulsars.

We briefly summarize the current status of pulsar search algorithms. Almost all pulsar searching is done using a variant of a power spectrum folding algorithm, which can be described as follows. First suppose for simplicity that the pulse period is constant. We square the Fourier transform of the data timestream $d(t)$ to obtain its power spectrum $P(\omega) = |\int d(t)e^{i\omega t} dt|^2$. A constant-period pulsar shows up in the power spectrum as a sharp peak at the pulse frequency ω_p , accompanied by a series of decaying sharp peaks at higher harmonics $\omega = (n\omega_p)_{n \geq 2}$. To sum these contributions, we “fold” the power spectrum, by computing $P_{\text{folded}}(\omega) = \sum_{n=1}^{n_{\text{max}}} W_n P(n\omega)$, for some fixed weighting W_n , and take the maximum (over ω) value of $P_{\text{folded}}(\omega)$ as our search statistic.

If the pulsar frequency is not assumed constant, say for example a constant-acceleration model is assumed, then one approach is to loop over trial accelerations, reparameterize the time coordinate to reduce to the constant-period case, and compute the folded power spectrum of the reparameterized timestream. Another approach is the Fourier-domain matched filter from [8].

Power spectrum folding is not optimal, even if the weighting W_n is matched to the power spectrum of the pulse profile. Because pulses are narrow in the time domain, the phases of the Fourier modes $\tilde{d}(n\omega_p)$ are not independent, and power spectrum folding throws away this phase information.

Given infinite computational power, an optimal coherent search could be implemented by brute force as follows. Assume a constant-acceleration model for simplicity. We loop over trial acceleration $\alpha = \ddot{\Phi}$, trial frequency $\omega = \dot{\Phi}$, and trial phase Φ . For each triple (α, ω, Φ) , we compute an intensity timestream $I_{\alpha\omega\Phi}(t)$, and compute its overlap $\hat{\mathcal{E}}(\alpha, \omega, \Phi) = \int dt d(t) I_{\alpha\omega\Phi}(t)$ with the data timestream $d(t)$. This is manifestly optimal but computationally slow. In this paper, we will show how to compute the transform $d(t) \rightarrow \hat{\mathcal{E}}(\alpha, \omega, \Phi)$ in a much faster way. In the case where there are no trial accelerations (constant-period search), the transform simply factors as a sequence of FFT’s (§III). The constant-acceleration search can be reduced to the constant-period case using a **fast recursive tree algorithm** (§IV).

The basic reason that pulsar searching is slow is that the size S of the search space, i.e. the number of independent pulsar models in the search, is a rapidly growing function of the timestream length T . For a constant-acceleration search, S is proportional to one power of T for $T \lesssim T_a$, and $S \propto T^3$ for $T \gtrsim T_a$, where T_a is the threshold timestream

length where the constant-period approximation breaks down. (For a pulsar with period P , period derivative \dot{P} , and duty cycle D , the threshold is parametrically $T_a \sim P \dot{P}^{-1/2} D^{1/2}$.) For large enough T , the constant-acceleration approximation will break down, and the search space size will grow as $S \propto T^6$ if a single parameter \dot{P} suffices, or even more rapidly if more parameters are needed.

In a strictly optimal coherent search, the overlap integral $\hat{\mathcal{E}} = \int dt d(t) I(t)$ is computed for every model $I(t)$ in the search space. A lower bound on the computational cost is $\mathcal{O}(S)$, where S is the size of the search space, since $\mathcal{O}(S)$ values of $\hat{\mathcal{E}}$ must be examined. Our recursive tree algorithm in §IV saturates this bound, but for large timestreams, S may still be too large for the optimal search to be practical. Therefore, we want to consider alternative search algorithms which trade off optimality for speed.

The main idea of this paper is a proposal for such an algorithm, *semicoherent search* (§VI). Our algorithm divides the data into chunks of size $T_c \ll T$, runs a coherent search in each chunk, and combines these subsearches using a procedure which keeps as much phase information as possible. The semicoherent search has an interpretation as the optimal search algorithm for a phase model whose acceleration $\ddot{\Phi}$ is allowed to wander over a small range. Even though the size of this search space is formally exponential in T , we will show that its optimal search statistic satisfies recursion relations which permit evaluation in $\mathcal{O}(T)$ time. Intuitively, letting $\ddot{\Phi}$ wander “fuzzes out” the search space so that there is a finite resolution to the phase information which must be retained, reducing computational cost.

Finally, we define a hierarchical search (§VII), which is simply a sequence of semicoherent searches with increasing coherence time T_c . The most significant peaks from each search are passed to the next search, as parameter ranges to be searched more optimally. As T_c increases, the semicoherent search becomes fully coherent. Thus, if a pulsar is above a certain signal-to-noise threshold, then it will show up as a rare peak in each level of the search, and the hierarchical search will converge to the complete phase model of the pulsar, as if a coherent search had been run. Of course, the key question is how this signal-to-noise threshold compares to a full coherent search, or to other algorithms (such as power spectrum stacking).

The main result of this paper is in §VIII, Fig. 1. We study the efficiency of the hierarchical search in Monte Carlo simulations. Remarkably, we find that the hierarchical search is nearly optimal for $T \lesssim 64T_c$. In other words, given a timestream which is 64 times larger than the longest timescale which can be searched coherently, we can still do a near-optimal search. Since the hierarchical search has cost $\mathcal{O}(T)$ and the coherent search has cost $\mathcal{O}(T^3)$, this should save a factor $64^2 = 4096$ in computing time. For $T \gtrsim 64T_c$, the hierarchical search is suboptimal, but the suboptimality grows slowly, and the hierarchical search is closer to optimal than power spectrum stacking.

We note a few caveats. First, our current code is a reference implementation and is not very well optimized. Therefore, in this paper we will make rough estimates of computational cost, rather than giving hard timings. We plan to improve this in follow-up work.

Second, we assume a fixed pulse profile throughout the analysis (a von Mises profile with duty cycle $D = 0.1$). In real data, we would need to add an outer loop over trial duty cycles. Note that in a coherent search, it suffices to run the search at the smallest duty cycle, then obtain larger duty cycles by smoothing the output $\hat{\mathcal{E}}(\alpha, \omega, \Phi)$. This procedure works because coherent search is a linear operation $d(t) \rightarrow \hat{\mathcal{E}}(\alpha, \omega, \Phi)$, but since the semicoherent search is nonlinear, we would need to rerun the search for each trial duty cycle.

Finally, throughout this paper, when we refer to a “timestream”, we mean a dedispersed intensity time series $I(t)$ at a fixed trial sky location and trial dispersion measure (DM). The reader should keep in mind that the total computational cost of processing a survey is larger by $(N_{\text{sky}} N_{\text{DM}})$, where N_{sky} is the number of sky pointings and N_{DM} is the number of trial DM’s. A full survey might have $N_{\text{sky}} \sim 10^4$ pointings, and $N_{\text{DM}} \sim 1300$ or $N_{\text{DM}} \sim 50000$ for a slow pulsar or millisecond pulsar search respectively. These N_{DM} values were derived assuming that the maximum DM of the search is $200 \text{ cm}^{-3} \text{ pc}$, the spacing between trial DM’s corresponds to a 1-sample delay across the full band, the survey has 100 MHz bandwidth at central frequency 400 MHz, and the sampling rate is $t_{\text{samp}} = 2 \text{ ms}$ for slow pulsars or $t_{\text{samp}} = 50 \text{ } \mu\text{s}$ for millisecond pulsars.

II. PRELIMINARIES

We model a pulsar by its phase model $\Phi(t)$ and pulse profile $\rho(\phi)$. The phase model maps time t to a dimensionless pulse phase $\Phi(t)$ such that the peak intensity occurs when $\Phi(t)$ is a multiple of 2π . In this paper, the most general

phase model we will consider is the constant-acceleration model, defined by

$$\Phi(t) = \Phi_0 + \omega_0 t + \frac{1}{2} \alpha t^2 \quad (1)$$

with parameters (t, ω_0, α) . The period P and its derivative \dot{P} are given by $(P, \dot{P}) = (2\pi\omega^{-1}, -2\pi\alpha\omega^{-2})$.

The pulse profile $\rho(\phi)$ gives the pulse intensity as a function of pulse phase ϕ . Throughout this paper, we will assume the von Mises profile

$$\rho(\phi) = e^{\kappa(\cos \phi - 1)} = e^{-2\kappa \sin^2(\phi/2)} \quad (2)$$

where κ is a concentration parameter. Alternatively, we can parameterize the von Mises profile by its duty cycle D , which we define to be the full width at half maximum (FWHM) of the pulse, divided by the pulse period. The parameters κ and D are related by $\kappa = (\log 2)/(2 \sin^2(\pi D/2))$.

The intensity timestream $I(t)$ of the pulsar is simply the composite function $I(t) = \rho(\Phi(t))$. We will assume that the timestream has been discretized with sample length t_s , so that the data is a sequence I_0, I_1, \dots, I_{N-1} given by boxcar-averaging $I(t)$:

$$I_k = \frac{1}{t_s} \int_{kt_s}^{(k+1)t_s} dt I(t) \quad (3)$$

We will assume that the noise is Gaussian and uncorrelated, i.e. the noise covariance is $\langle I_k I_l \rangle = \eta^2 t_s^{-1} \delta_{kl}$. Here, η is a noise parameter with units intensity-(time) $^{1/2}$.

We write \bar{I}_k for the discretized signal normalized to signal-to-noise 1, i.e. obeying normalization condition

$$\eta^{-2} t_s \sum_{k=0}^{N-1} (\bar{I}_k)^2 = 1 \quad (4)$$

Note that the normalization of \bar{I}_k implicitly depends on the number of timestream samples N .

Given data realization d_k and fixed pulsar model I_k , the optimal statistic for detecting the pulsar is the coherent detection statistic $\hat{\mathcal{E}}$ defined by:

$$\hat{\mathcal{E}} = \eta^{-2} t_s \sum_{k=0}^{N-1} d_k \bar{I}_k \quad (5)$$

By “coherent”, we mean that all phase information is used. We have normalized $\hat{\mathcal{E}}$ so that its numerical value is the detection significance in sigmas. Equivalently, $\hat{\mathcal{E}}$ is equal to $(\Delta\chi^2)$, the improvement in χ^2 after subtracting a best-fit multiple of the pulsar model \bar{I}_k . This normalization is convenient, but note that if two timestreams of lengths N_1, N_2 are combined, the rule for combining $\hat{\mathcal{E}}$ -values is:

$$\hat{\mathcal{E}} = \frac{N_1^{1/2}}{(N_1 + N_2)^{1/2}} \hat{\mathcal{E}}_1 + \frac{N_2^{1/2}}{(N_1 + N_2)^{1/2}} \hat{\mathcal{E}}_2 \quad (6)$$

III. FAST CONSTANT-PERIOD SEARCH

In this section we will consider the simplest possible search space: a pulsar with constant frequency, parameterized by a frequency ω and phase Φ . A brute-force optimal search algorithm for this search space would be to loop over a grid of trial parameters (ω, Φ) , and compute the coherent statistic $\hat{\mathcal{E}}$ defined in Eq. (5), which becomes a function $\hat{\mathcal{E}}(\omega, \Phi)$. In this section we will show that the brute-force search can be computed in a mathematically equivalent but faster way. If we view the coherent search as a transform $d(t) \rightarrow \hat{\mathcal{E}}(\omega, \Phi)$, then the transform can be factored as a sequence of FFT's.

The phase model for this search is:

$$\Phi(t) = \Phi_c + \omega_c \left(t - \frac{T}{2} \right) \quad (0 \leq t \leq T) \quad (7)$$

where $T = Nt_s$ is the total timestream length, ω_c is the angular pulse frequency, and Φ_c is the pulse phase at the timestream center $t = T/2$. Note that we have taken the model parameters to be frequency and central phase Φ_c , whereas previously in Eq. (1) we used the initial phase $\Phi_0 = \Phi_c - \omega T/2$ instead of Φ_c . This change of variables will be convenient for reasons to be explained shortly.

We Fourier transform the pulse profile $\rho(\phi)$:

$$\rho(\phi) = \sum_n \rho_n e^{in\phi} \quad (8)$$

where the sum runs over positive and negative n . We then write the signal timestream $I(t)$ as

$$\begin{aligned} I(t) &= \rho(\Phi(t)) \\ &= \rho\left(\Phi_c + \omega_c \left(t - \frac{T}{2}\right)\right) \\ &= \sum_n \rho_n e^{in\Phi_c} e^{in\omega_c(t-T/2)} \end{aligned} \quad (9)$$

The discretized timestream I_k is obtained from the continuous timestream $I(t)$ by boxcar averaging. Equivalently, we can convolve $I(t)$ with a length- t_s boxcar, then evaluate at the sample center $t = (k + 1/2)t_s$. The convolution can be implemented by multiplying each Fourier mode of $I(t)$ by $j_0(\omega t_s/2)$, the Fourier transform of a length- t_s boxcar. Therefore, starting from Eq. (9) the discretized timestream can be written:

$$I_k = \sum_n \rho_n j_0\left(\frac{n\omega_c t_s}{2}\right) e^{in\Phi_c} e^{in\omega_c(kt_s + t_s/2 - T/2)} \quad (10)$$

where $j_0(x) = (\sin x)/x$.

The normalized timestream \bar{I}_k defined in the previous section is given by $\bar{I}_k = A^{-1/2} I_k$, where the normalization A is given by:

$$A = \eta^{-2} t_s \sum_k (I_k)^2 \quad (11)$$

We next derive an approximate formula for A , in the limit where the timestream length T is large compared to the pulse period. We approximate the sum, which has the schematic form $\eta^{-2} t_s \sum_k I(k t_s)^2$, by the integral $\eta^{-2} \int dt I(t)^2$, and plug in Eq. (10) to obtain:

$$A \approx \eta^{-2} \int dt \left[\sum_n \rho_n j_0\left(\frac{n\omega_c t_s}{2}\right) e^{in\Phi_c} e^{in\omega_c(t + t_s/2 - T/2)} \right]^2 \quad (12)$$

We now expand the square, making the approximation $\int dt e^{im\omega_c t} e^{in\omega_c t} \approx T \delta_{m,-n}$. This gives:

$$A(\omega_c) \approx \eta^{-2} T \sum_n |\rho_n|^2 j_0\left(\frac{n\omega_c t_s}{2}\right)^2 \quad (13)$$

where we have written $A(\omega_c)$ on the LHS to emphasize that it depends on ω_c but not Φ_c . We can use this formula to get the correct normalization for \bar{I}_k , starting from an arbitrary normalization for the pulse profile ρ_n .

Now we obtain a formula for the search statistic $\hat{\mathcal{E}}$, by plugging Eq. (10) into the definition (5) of $\hat{\mathcal{E}}$:

$$\hat{\mathcal{E}}(\omega_c, \Phi_c) = A(\omega_c)^{-1/2} \eta^{-2} t_s \sum_k d_k \sum_n \rho_n j_0\left(\frac{n\omega_c t_s}{2}\right) e^{in\Phi_c} e^{in\omega_c(kt_s + t_s/2 - T/2)} \quad (14)$$

To simplify this, we define the Fourier transform of the data realization d_k by:

$$\tilde{d}(\omega) = t_s \sum_k d_k e^{i\omega(kt_s + t_s/2 - T/2)} \quad (15)$$

Note that this definition contains an extra phase $e^{i\omega(t_s/2 - T/2)}$ relative to the usual Fourier transform. Then $\hat{\mathcal{E}}(\omega_c, \Phi_c)$ can be written

$$\hat{\mathcal{E}}(\omega_c, \Phi_c) = \frac{1}{A(\omega_c)^{1/2} \eta^2} \sum_n \rho_n j_0\left(\frac{n\omega_c t_s}{2}\right) \tilde{d}(n\omega_c) e^{in\Phi_c} \quad (16)$$

This formula can be used to give a fast algorithm for evaluating $\hat{\mathcal{E}}(\omega_c, \Phi_c)$. First, we precompute $\tilde{d}(\omega)$ on a grid of ω -values, by zero-padding the timestream and taking an FFT. The purpose of the zero-padding is to make the ω -sampling narrow enough that $\tilde{d}(\omega)$ can be evaluated at an arbitrary frequency by interpolation. We find that zero-padding by a factor two is sufficient (Appendix A). Similarly, we precompute ρ_n , and populate an interpolation table with precomputed values of $A(\omega)$ using Eq. (13).

Now to evaluate $\hat{\mathcal{E}}(\omega_c, \Phi_c)$ on a grid of trial (ω_c, Φ_c) values, we evaluate Eq. (16) by interpolating the factors which appear on the RHS. To do the n -sum efficiently, we use an FFT from the variable n to variable Φ_c .

This concludes our fast FFT-based algorithm for constant-frequency search, but there are a few details which deserve elaboration. We have chosen to use the central phase Φ_c as a model parameter in Eq. (7), rather than the initial phase Φ_0 . This resulted in an phase shift $e^{-i\omega(T/2)}$ in the definition (15) of $\tilde{d}(\omega)$. Empirically, we find that this choice results in a more well-behaved $\tilde{d}(\omega)$ interpolation. This can also be understood formally by noticing that the phase $e^{i\omega(kt_s + t_s/2 - T/2)}$ appearing in Eq. (15) has fewer wraparounds (as ω is varied) with the factor $e^{-i\omega T/2}$ than without.

There are a hidden parameters in our algorithm: the amount of zero-padding used to compute the timestream FFT, and the spacings of the trial parameters Φ_c and ω_c . Strictly speaking, our algorithm is only optimal in the limit of large zero-padding and small trial spacings. In practice we choose parameter values which are a compromise between optimality and computational cost. In Appendix A we give a formal procedure for making these choices and recommend some default values.

In real data analysis, the timestream d_k must be “detrended” or high-pass filtered before being searched for pulsars. It would be equivalent to leave the data d_k unfiltered, but apply high-pass filtering to the pulsar signal \bar{I}_k before computing the search statistic $\hat{\mathcal{E}}$. Therefore, we can account for the effect of detrending by simply subtracting the mean from the pulse profile $\rho(\phi)$, or equivalently setting the Fourier mode ρ_0 to zero. All results in this paper include this detrending correction, but it turns out to make little difference.

The computational cost of our fast algorithm can be estimated as follows. The initial FFT has cost $\mathcal{O}(N \log N)$, where N is the number of time samples, and the $\hat{\mathcal{E}}$ calculation has cost $\mathcal{O}(N_\omega N_\Phi \log N_\Phi)$, where N_ω, N_Φ are the number of trial ω and Φ values needed. It is not hard to see that $N_\Phi = \mathcal{O}(D^{-1})$ and $N_\omega = \mathcal{O}(\omega_{\max} T D^{-1})$, where D is the duty cycle and ω_{\max} is the maximum value of ω which is searched. We will assume that the data has been sampled so that the sample length t_s is comparable to the minimum pulse width in the search, or equivalently $N = \mathcal{O}(\omega_{\max} T D^{-1})$. Putting all of this together, we can write the total cost as $\mathcal{O}(N \log N + N D^{-1} \log D^{-1})$.

This cost can be compared to the power spectrum folding algorithm described in the introduction, which has cost $\mathcal{O}(N \log N)$. The computational costs will typically be comparable, but our coherent algorithm will be significantly slower in the limit of low duty cycle. On the other hand, this is also the limit where power spectrum folding becomes significantly suboptimal, so one could argue that it is always a good idea to use the fast coherent search instead of power spectrum folding.

IV. TREE ALGORITHM FOR CONSTANT-ACCELERATION SEARCH

We now consider a more complex case: a constant-acceleration search. In this case, the output of the coherent search will be a function of three variables $\hat{\mathcal{E}}(\alpha, \omega, \phi)$, where $\alpha = \ddot{\Phi}$ is an acceleration parameter. We will give a fast algorithm for evaluating $\hat{\mathcal{E}}$ on a grid of trial parameters.

The algorithm is recursive and based on the following idea. We divide the time interval $[0, T]$ into two subintervals $[0, T/2]$ and $[T/2, T]$. At any point (α, ω, ϕ) , the statistic $\hat{\mathcal{E}}$ will be a sum of contributions $\hat{\mathcal{E}} = \hat{\mathcal{E}}_1 + \hat{\mathcal{E}}_2$. The number of grid points needed to fully represent $\hat{\mathcal{E}}_1$ and $\hat{\mathcal{E}}_2$ will smaller by a factor $(1/8)$, since the number of trial accelerations scales as T^2 and the number of trial frequencies scales as T . Therefore we can compute $\hat{\mathcal{E}}_1$ and $\hat{\mathcal{E}}_2$ on coarser grids, and interpolate to a finer grid when we sum them to obtain $\hat{\mathcal{E}}$. The values of $\hat{\mathcal{E}}_1$ and $\hat{\mathcal{E}}_2$ are computed recursively using the same method, further subdividing the time range. After enough subdivisions, the timestream will be short

enough that no trial accelerations are needed, and the fast constant-period search from the previous section can be used to compute $\hat{\mathcal{E}}$, ending the recursion.

We parameterize the phase model as:

$$\Phi(t) = \bar{\Phi} + \omega_c \left(t - \frac{T}{2} \right) + \frac{1}{2} \alpha \left[\left(t - \frac{T}{2} \right)^2 - \frac{T^2}{12} \right] \quad (17)$$

over the range $0 \leq t \leq T$. Note that we have changed variables from the parameterization $(\alpha, \omega_0, \Phi_0)$ given previously in Eq. (1) to the parameters $(\alpha, \omega_c, \bar{\Phi})$. The parameter ω_c is the derivative $d\Phi/dt$ evaluated at the central time $t = T/2$, and $\bar{\Phi}$ is the mean value of Φ over the range $0 \leq t \leq T$, as suggested by the notation. An analogous change of variables was made in the last section (Eq. (7)), in order to make an interpolation better behaved. The motivation here is similar and will be described shortly.

Let $\hat{\mathcal{E}}_1, \hat{\mathcal{E}}_2$ denote the search statistic $\hat{\mathcal{E}}$ restricted to the subinterval $[0, T/2]$ or $[T/2, T]$. When we write $\hat{\mathcal{E}}_1(\alpha, \omega_c, \bar{\Phi})$, the arguments $(\omega_c, \bar{\Phi})$ are always defined relative to the subinterval $[0, T/2]$, not the larger interval $[0, T]$ (and likewise for $\hat{\mathcal{E}}_2$). In this notation, a short calculation gives the recursion relating $\hat{\mathcal{E}}$ to $\hat{\mathcal{E}}_1, \hat{\mathcal{E}}_2$:

$$\hat{\mathcal{E}}(\alpha, \omega_c, \bar{\Phi}) = \frac{1}{\sqrt{2}} \hat{\mathcal{E}}_1 \left(\alpha, \omega_c - \frac{\alpha T}{4}, \bar{\Phi} - \frac{\omega_c T}{4} \right) + \frac{1}{\sqrt{2}} \hat{\mathcal{E}}_2 \left(\alpha, \omega_c + \frac{\alpha T}{4}, \bar{\Phi} + \frac{\omega_c T}{4} \right) \quad (18)$$

Now suppose that the search statistics $\hat{\mathcal{E}}_1, \hat{\mathcal{E}}_2$ have been precomputed on a complete grid of trial $(\alpha, \omega_c, \bar{\Phi})$ values. We use the recursion relation (18) to evaluate $\hat{\mathcal{E}}$ on a complete grid, using interpolation to obtain values of $\hat{\mathcal{E}}_1, \hat{\mathcal{E}}_2$ off-grid. The precomputation of $\hat{\mathcal{E}}_1$ and $\hat{\mathcal{E}}_2$ is done by applying the same idea recursively, initializing $\hat{\mathcal{E}}_1$ from tables $\hat{\mathcal{E}}_{11}$ and $\hat{\mathcal{E}}_{12}$ which correspond to intervals $[0, T/4]$ and $[T/4, T/2]$, and so on. After enough subdivisions, the subinterval size $(T/2^N)$ is small enough that the acceleration α is negligible, and we can approximate $\hat{\mathcal{E}}(\alpha, \omega_c, \bar{\Phi}) \approx \hat{\mathcal{E}}(0, \omega_c, \bar{\Phi})$. At this point, the grid of $\hat{\mathcal{E}}$ values can be computed using the constant-period algorithm from the previous section, which ends the recursion.

The computational cost of the search can be computed as follows. If we write the total number of $\hat{\mathcal{E}}$ evaluations needed at the top level of the search as $N_\alpha N_\omega N_\Phi$, then the total number of $\hat{\mathcal{E}}_1$ and $\hat{\mathcal{E}}_2$ evaluations needed at the second level is $2(N_\alpha/4)(N_\omega/2)N_\Phi = N_\alpha N_\omega N_\Phi/4$. Similarly, the total number of evaluations needed at the third level of the search is $N_\alpha N_\omega N_\Phi/16$, and so on. Summing over all levels of the search, the total computational cost is $\mathcal{O}(N_\alpha N_\omega N_\Phi + N \log N + ND^{-1} \log D^{-1})$, where the first term is the sum of a convergent geometric series, and the second and third terms are the total cost of all the constant-period searches at the bottom level. Provided that $N_\alpha \gtrsim \log N$ and $N_\alpha \gtrsim D^{-1} \log D^{-1}$, which will be the case for all large constant-acceleration searches, the first term dominates and we can write the total cost as $\mathcal{O}(N_\alpha N_\omega N_\Phi)$. Remarkably, we see that the computational cost of the tree search is $\mathcal{O}(1)$ per model in the search space!

The tree algorithm contains hidden parameters: the grid spacings in the parameters $(\alpha, \omega_c, \bar{\Phi})$ used to construct interpolation tables at each level of the search, and the threshold for switching to the bottom-level constant-period search. On a related note, the reader may wonder whether there is any suboptimality introduced cumulatively by the chain of interpolations in the tree search. In Appendix A, we show that it is straightforward to choose parameters so that the search is as close to optimal as desired. We give our default parameter choices, and show that they produce a search which is $> 90\%$ optimal.

A very interesting aspect of the tree algorithm is that it should generalize straightforwardly to search spaces more complex than the constant-acceleration search, for example a polynomial search of degree 3 or 4. As the search recurses, trial parameter spacings can be decreased, and the polynomial degree can also be decreased. We defer an exploration of more complex search spaces to future work.

The parameterization $(\alpha, \omega_c, \bar{\Phi})$ was chosen in Eq. (17) because we find empirically that the interpolation $\hat{\mathcal{E}}(\alpha, \omega_c, \bar{\Phi})$ is better behaved, than a simpler parameterization such as $(\alpha, \omega_0, \Phi_0)$. In fact, the parameterization in Eq. (17) is just the Legendre polynomial expansion truncated at degree 2 and rescaled to the interval $[0, T]$. Since Legendre polynomials are orthogonal, varying their coefficients produces an uncorrelated effect on $\hat{\mathcal{E}}$, which leads to a more efficient interpolation. This way of thinking about the parameterization makes the generalization to higher-degree polynomial models transparent.

So far, we have discussed searches which are strictly optimal, in the sense that the optimal statistic $\hat{\mathcal{E}}$ is evaluated for every model in the search space. The tree algorithm from this section has computational cost $\mathcal{O}(S)$, where $S = N_\alpha N_\omega N_\Phi$ is size of the search space, i.e. the number of independent phase models.

It may seem plausible that this is a lower bound on the computational cost of any optimal search, by the following argument. Even if $\hat{\mathcal{E}}$ had been precomputed for every model and stored in memory, we still need to inspect S values of $\hat{\mathcal{E}}$ in order to perform the search, and the cost of simply reading these values from memory is $\mathcal{O}(S)$. Surprisingly, this simple argument is incorrect! As we will see in §VI, there are examples of “fuzzy” search spaces which can be optimally searched with computational cost $\mathcal{O}(\log S)$, where S is the size of the space. Although a constant-acceleration search is not an example of a fuzzy search space, there is a sense in which it can be approximated by one. This will lead us to the notion of a semicoherent search, the main idea of this paper.

V. THE $\hat{\mathcal{H}}$ -STATISTIC

In this section, we study the following question. Consider a discrete search space which consists of S distinct pulsar models $I_k^{(1)}, \dots, I_k^{(S)}$, and suppose we have evaluated $\hat{\mathcal{E}}$ for each model, obtaining S values $\hat{\mathcal{E}}_1, \dots, \hat{\mathcal{E}}_S$. If we want to compress these S numbers into a single number which represents the result of the search, what should we do? Should we use the max-statistic $\max(\hat{\mathcal{E}}_1, \dots, \hat{\mathcal{E}}_S)$, the χ^2 -like statistic $\sum_{s=1}^S (\hat{\mathcal{E}}_s)^2$, or something else?

Let us interpret this question in the language of frequentist hypothesis testing. We would like to compare two hypotheses, a null hypothesis that the data is pure noise, and an alternative hypothesis that the data is noise plus a pulsar signal. To specify the alternative hypothesis precisely, let us suppose that $r\bar{I}_k^{(s)}$ is added to the noise, where r is a fiducial signal-to-noise in sigmas, and $s = 1, \dots, S$ is a random pulsar model.

By the Neyman-Pearson lemma, the optimal statistic for distinguishing these hypotheses is the likelihood ratio $(\mathcal{L}_1/\mathcal{L}_0)$, where \mathcal{L}_0 and \mathcal{L}_1 are the likelihoods of the data given the null and alternative hypotheses. To write these likelihoods compactly, we introduce a dot product notation. Given two discretized timestreams x_k, x'_k , we define their dot product by:

$$x \cdot x' = \eta^{-2} t_s \sum_k x_k x'_k \quad (19)$$

In this notation, the estimator $\hat{\mathcal{E}}$ is defined by $\hat{\mathcal{E}} = d \cdot \bar{I}$, and the normalization of \bar{I} defined in Eq. (4) is $\bar{I} \cdot \bar{I} = 1$. Neglecting overall constants, the likelihoods \mathcal{L}_0 and \mathcal{L}_1 are given by

$$\mathcal{L}_0 \propto e^{-(d \cdot d)/2} \quad \mathcal{L}_1 \propto \frac{1}{S} \sum_{s=1}^S e^{-(d - r\bar{I}^{(s)}) \cdot (d - r\bar{I}^{(s)})/2} \quad (20)$$

After a little algebra, the likelihood ratio can be written as $(\mathcal{L}_1/\mathcal{L}_0) \propto e^{r\hat{\mathcal{H}}}$, where we define the $\hat{\mathcal{H}}$ -statistic by:

$$\hat{\mathcal{H}} = \frac{1}{r} \log \frac{1}{S} \sum_{s=1}^S e^{r\hat{\mathcal{E}}_s} \quad (21)$$

We can think of $\hat{\mathcal{H}}$ as a version of $\hat{\mathcal{E}}$ which has been coarse-grained over a model space of size S . Note that the definition of $\hat{\mathcal{H}}$ depends both on the model space and the fiducial signal-to-noise r . The statistic $\hat{\mathcal{H}}$ is only strictly optimal if r is equal to the true signal-to-noise of the pulsar. In practice, we set r equal to the detection threshold of the search. It makes sense to optimize the search for pulsars near threshold, since pulsars with signal-to-noise significantly above threshold will still be detected if the statistic is slightly suboptimal, and pulsars significantly below threshold will never be detected.

Note that $\hat{\mathcal{H}}$ is a likelihood ratio test in disguise, since it is just a reparameterization of $(\mathcal{L}_1/\mathcal{L}_0)$. This reparameterization is convenient because the value of $\hat{\mathcal{H}}$ on a timestream containing a bright pulsar (i.e. significantly above threshold) is simply its coherent signal-to-noise in sigmas. We also note that $\hat{\mathcal{H}} = \hat{\mathcal{E}}$ if the search space size $S = 1$.

We have now answered the question from the beginning of this section: the statistic $\hat{\mathcal{H}}$ defined in Eq. (21) compresses the values $\hat{\mathcal{E}}_1, \dots, \hat{\mathcal{E}}_S$ into a single optimal statistic. As an aside, we note that $\hat{\mathcal{H}}$ interpolates between the max-statistic $\max(\hat{\mathcal{E}}_1, \dots, \hat{\mathcal{E}}_S)$ and the χ^2 -like statistic $\sum_{s=1}^S (\hat{\mathcal{E}}_s)^2$ in the limits $r \rightarrow \infty$ and $r \rightarrow 0$ respectively. The first statement is easy to show, and the second statement follows by Taylor expanding

$$\hat{\mathcal{H}} = \frac{1}{S} \sum_s \hat{\mathcal{E}}_s - \frac{r}{2S^2} \left(\sum_s \hat{\mathcal{E}}_s \right)^2 + \frac{r}{2S} \sum_s (\hat{\mathcal{E}}_s)^2 + \mathcal{O}(r^2) \quad (22)$$

We write the quantity $\sum_s \hat{\mathcal{E}}_s$ as $d \cdot (\sum_s \bar{I}_s)$, and note that $\sum_s \bar{I}_s$ vanishes if the pulsar profile is detrended and the search space is invariant under the phase shift symmetry $\Phi \rightarrow \Phi + \Delta\Phi$. Therefore the first two terms in the Taylor series (22) vanish, and the leading term as $r \rightarrow 0$ is the third term, which is proportional to the χ^2 -like statistic $\sum_{s=1}^S (\hat{\mathcal{E}}_s)^2$.

VI. SEMICOHERENT SEARCH

In this section, we introduce the notion of a semicoherent search. Suppose we are interested in a constant-acceleration search over acceleration range $[-\alpha_0, \alpha_0]$. We divide the timestream into $N_c \gg 1$ chunks, where the chunk duration $T_c = T/N_c$ is assumed comparable to the acceleration timescale $T_a = (D/\alpha_0)^{1/2}$. Under these assumptions, a constant-period search would suffice if we restrict the search to a single length- T_c chunk, but we would need to use multiple trial α -values to search larger chunk sizes.

Let us imagine that it is computationally feasible to do a constant-period search in each length- T_c chunk, but we do not have enough computing power to analyze larger chunk sizes. How can we stitch together the per-chunk search statistics $\hat{\mathcal{E}}_i(\omega_c, \Phi_c)$, where $i = 1, \dots, N_c$, in order to form the best possible search statistic? Of course, this will be a suboptimal statistic, since we are assuming that the fully coherent search is not affordable, but we are asking for the *least suboptimal* way of assembling coherent searches computed on the timescale $T_c \ll T$.

Informally, the idea of the semicoherent search is to compute the $\hat{\mathcal{H}}$ -statistic for a “fuzzy” search space consisting of all phase models whose acceleration $\alpha = \ddot{\Phi}$ is allowed to wander over the range $[-\alpha_0, \alpha_0]$. In each length- T_c chunk, the constant-period approximation will be valid. The $\hat{\mathcal{H}}$ -statistic for the fuzzy search space will combine the per-chunk $\hat{\mathcal{E}}$ values in a way which keeps as much phase consistency between the chunks as possible.

Formally, we introduce the following search space. We consider phase models which have the quadratic form $\Phi(t) = (1/2)\alpha t^2 + \omega t + \Phi_0$ in each chunk, but we allow the acceleration $\alpha = \ddot{\Phi}$ to change discontinuously at chunk boundaries. Specifically, we assume that $\alpha = \pm\alpha_0$, where the sign is independent in each chunk. We require the phase Φ and its derivative $\dot{\Phi}$ to be continuous across chunk boundaries. Thus we can parameterize the search space by initial phase Φ_0 , an initial frequency ω_0 , and a sequence of signs $s_i = \pm 1$, where $1 \leq i \leq N_c$. This data suffices to determine the phase model $\Phi(t)$ for all times $0 \leq t \leq T$.

The size of this search space grows exponentially with T , so it would be natural to assume that the computational cost of searching it is also exponential. However, we will now show that there is a recursion relation which allows the $\hat{\mathcal{H}}$ -statistic for this search space to be computed in linear time.

Given a pulsar in the search space, the coherent search statistic $\hat{\mathcal{E}}$ will be a sum over contributions from each chunk. Schematically, $\hat{\mathcal{E}} = (\hat{\mathcal{E}}_1 + \dots + \hat{\mathcal{E}}_{N_c})/\sqrt{N_c}$. Now suppose we define an $\hat{\mathcal{H}}$ -statistic by summing over all 2^{N_c} phase models with initial phase and frequency (Φ_0, ω_0) . We can write $\hat{\mathcal{H}}$ as an iterated sum:

$$\hat{\mathcal{H}}(\omega_0, \Phi_0) = \frac{1}{r} \log \sum_{s_1} \dots \sum_{s_{N_c}} \exp \left(\frac{r}{\sqrt{N_c}} \hat{\mathcal{E}}_1 + \dots + \frac{r}{\sqrt{N_c}} \hat{\mathcal{E}}_{N_c} \right) \quad (23)$$

where it is understood that each $\hat{\mathcal{E}}_i$ on the RHS is evaluated on an appropriate set of model parameters, which depend on (ω_0, Φ_0) and the signs $s_j = \pm 1$ with $j \leq i$.

Note that the definition in Eq. (23) partially coarse-grains the search space: we have coarse-grained over the signs s_i , but not the phase Φ_0 or frequency ω_0 .

We now write a recursion relation which allows $\hat{\mathcal{H}}$ to be computed efficiently. Given a chunk index $0 \leq j < N_c$, and a frequency and phase (ω_j, Φ_j) defined at the intermediate time $t = jT_c$, we define the partially summed statistic $\hat{\mathcal{H}}_j(\omega_j, \Phi_j)$ by:

$$\hat{\mathcal{H}}_j(\omega_j, \Phi_j) = \frac{1}{r} \log \sum_{s_{j+1}} \dots \sum_{s_{N_c}} \exp \left(\frac{r}{\sqrt{N_c}} \hat{\mathcal{E}}_{j+1} + \dots + \frac{r}{\sqrt{N_c}} \hat{\mathcal{E}}_{N_c} \right) \quad (24)$$

We have coarse-grained over the 2^{N_c-j} choices of sign which connect the intermediate time jT_c and the final time $T = N_c T_c$.

There is a recursion relation relating $\hat{\mathcal{H}}_j$ and $\hat{\mathcal{H}}_{j+1}$. To see this, note that if we fix the index s_{j+1} in the outermost sum in Eq. (24), then the iterated sum which remains is the same sum which appears in the definition of $\hat{\mathcal{H}}_{j+1}$. More

precisely, a short calculation gives the recursion:

$$\hat{\mathcal{H}}_j(\omega_j, \Phi_j) = \frac{1}{r} \log \sum_{s=\pm 1} \exp \left[\frac{r}{\sqrt{N_c}} \hat{\mathcal{E}}_j \left(\omega_j + s \frac{\alpha_0 T_c}{2}, \Phi_j + \frac{\omega_j T_c}{2} + s \frac{\alpha_0 T_c^2}{6} \right) + r \hat{\mathcal{H}}_{j+1} \left(\omega_j + s \alpha_0 T_c, \Phi_j + s \frac{\alpha_0 T_c^2}{2} \right) \right] \quad (25)$$

This recursion can be used to compute $\hat{\mathcal{H}}_j$ in the order $j = N_c, \dots, 0$. The initial condition for the recursion is $\hat{\mathcal{H}}_{N_c}(\omega, \Phi) = 0$. The output of the semicoherent search is a 2D array of values $\hat{\mathcal{H}}_0(\omega_0, \Phi_0)$ with shape (N_ω, N_Φ) , where N_ω is the number of trial frequencies associated with timestream length T_c (not length $T = N_c T_c$). The computational cost of iterating the recursion is $\mathcal{O}(N_c N_\omega N_\Phi)$. For a fixed chunk size T_c , the cost grows as one power of T , in contrast to a coherent search, whose cost would grow as T^3 .

VII. HIERARCHICAL SEARCH

First, we introduce a small generalization of the semicoherent search. In the previous section, we defined a single acceleration bin $[-\alpha_0, \alpha_0]$, and allowed $\ddot{\Phi}$ to wander over this range. As a generalization, suppose we cover the acceleration range of the search with N_α acceleration bins, with bin width $(\Delta\alpha)$. We define a multi-bin semicoherent search, by modifying the single-bin construction above as follows:

- We choose the chunk size T_c comparable to the timescale $(D/\Delta\alpha)^{1/2}$ where the trial acceleration spacing of a coherent search is $(\Delta\alpha)$. This will be a longer timescale than in the single-bin case, where $T_c \sim T_a = (D/\alpha_0)^{1/2}$ is the timescale where no trial accelerations are needed at all.
- In each length- T_c chunk, we run a coherent search, which will have N_α trial accelerations, using the recursive tree algorithm from §IV. (In the single-bin case, we chose T_c so that no trial accelerations were necessary, and used the fast constant-frequency search algorithm from §III.)
- For each acceleration bin a , initial frequency ω_0 , and initial phase Φ_0 , we coarse-grain over all phase models whose acceleration $\ddot{\Phi}$ is allowed to wander over the bin a . The resulting $\hat{\mathcal{H}}$ -statistic $\hat{\mathcal{H}}(a, \omega_0, \Phi_0)$ can be computed using the recursion relation (25).
- The output of the semicoherent search is a 3D array of values $\hat{\mathcal{H}}(a, \omega_0, \Phi_0)$ with shape $(N_\alpha, N_\omega, N_\Phi)$. Here, N_α is the number of acceleration bins which were specified as an input parameter of the search, and N_ω is the number of trial frequencies associated with timestream length T_c . The computational cost is $\mathcal{O}(N_\alpha N_\omega N_\Phi)$.

We can think of the semicoherent search as being parameterized by an α -bin width $(\Delta\alpha)$, or by a coherence length T_c . The two will be related by $T_c \propto (D/\Delta\alpha)^{1/2}$, but the constant of proportionality is a parameter that can be optimized. In practice, we usually determine T_c from considerations of computational cost: we simply set it to the largest chunk size where we can afford to do a coherent search. The semicoherent search can be interpreted as a procedure for stitching together the results of coherent searches on the timescale $T_c \ll T$, to obtain a single search statistic $\hat{\mathcal{H}}$ which preserves as much phase information as possible.

One more definition. A “hierarchical search” is just a sequence of semicoherent searches with increasing values of T_c . (In our implementation, we take each T_c to be 4 times larger than the previous level of the hierarchy.) At each level of the hierarchy, the rare peaks in the output array $\hat{\mathcal{H}}(a, \omega_0, \Phi_0)$ are used to define input search ranges for the search at the next level of hierarchy. Because only a small fraction of the search space is passed on, the total computational cost in our implementation is dominated by the first level of the hierarchy.

The last level of the hierarchy is a coherent search, which only runs on a very small fraction of the full (α, ω, Φ) parameter space. If the timestream contains a bright pulsar, then it will produce a rare peak at every level of the hierarchy, and the hierarchical search will return a fit for the full phase model, as if a coherent search had been run. Of course, the key question is, what is the signal-to-noise threshold for the hierarchical search to work? In the next section, we study this question using Monte Carlo simulations.

VIII. MONTE CARLO SIMULATIONS

In this section, we run Monte Carlo simulations using the following fiducial survey parameters. We assume time sampling $t_{\text{samp}} = 50 \mu\text{sec}$ as appropriate for millisecond pulsars, and total timestream size 2^{28} samples, corresponding

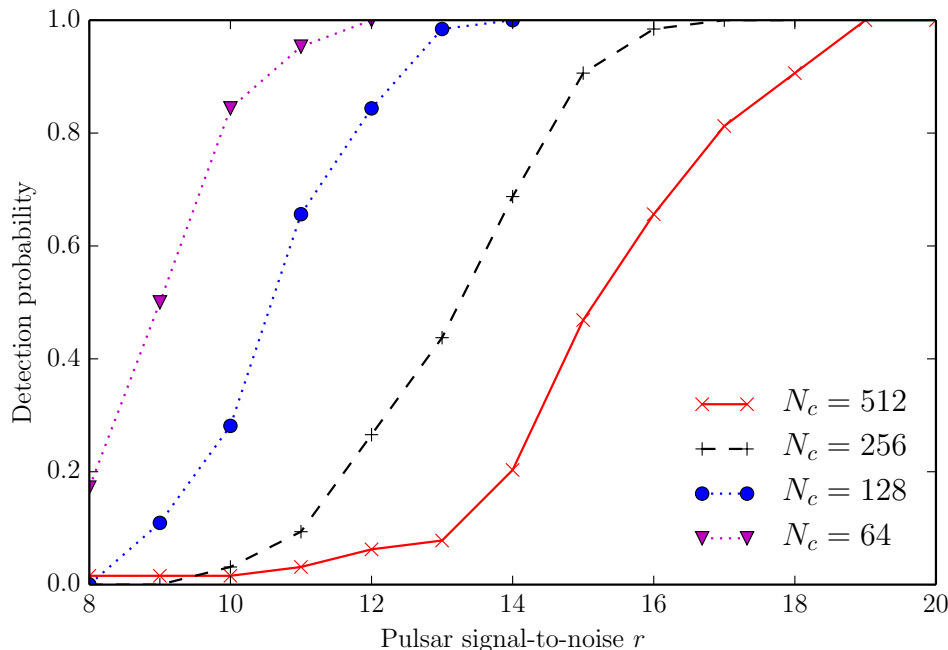


FIG. 1: Monte Carlo detection probability of hierarchical search, as a function of the number of coherent chunks N_c and the total signal-to-noise r of the simulated pulsar. In the text, we interpret these detection probabilities and show that the hierarchical search is close to optimal at $N_c = 64$. For larger N_c it is less suboptimal than incoherent power spectrum stacking, an alternative chunk-based method.

to integration time $T = 3.7$ hours. Our search assumes a von Mises profile with fixed duty cycle $D = 0.1$.

We search the frequency range $(\omega_{\min}, \omega_{\max}) = (200, 4000)$ Hz, corresponding to pulsar period P between 1.5 and 31 milliseconds. We search acceleration range $(\alpha_{\min}, \alpha_{\max}) = (-10^{-4}, 10^{-4}) \text{ sec}^{-2}$. For relativistic binary pulsars with $v/c = 10^{-3}$, this acceleration range covers orbital periods P_{orb} greater than 3.5 hours if $P = 31$ msec, or P_{orb} greater than 70 hours if $P = 1.5$ msec.¹

With these parameters, we estimate that a fully coherent search would have search space size $S = N_\alpha N_\omega N_\phi \approx 2 \times 10^{13}$. This represents the computational cost of a single sky pointing and trial DM. In a full millisecond pulsar survey, the cost would be larger by a factor $N_{\text{sky}} N_{\text{DM}} \sim 10^9$ as described in the introduction. Therefore, a full coherent search is not computationally feasible for a full survey with these parameters.

We run hierarchical searches, parameterized by the number of chunks N_c in the first level of the hierarchy. Since the first level dominates the computational cost, the cost is the same as N_c coherent searches with timestream length $T_c = T/N_c$. Since the cost of a coherent search is $\mathcal{O}(T_c^3)$, the hierarchical search is faster than a full coherent search by a factor N_c^2 . In each Monte Carlo simulation, we simulate a pulsar with random parameters, and say that the hierarchical search detects the pulsar if it converges to the full phase model of the pulsar at the last stage of the hierarchy. In Fig. 1, we show the detection probability of the search as a function of N_c and the signal-to-noise.

We emphasize that the total SNR on the x -axis is the total signal-to-noise of the pulsar summed over all N_c chunks. The hierarchical search succeeds in finding pulsars whose signal-to-noise per coherent chunk is very small. For example, a search with $r = 18$ and $N_c = 512$ succeeds 90% of the time, even though the signal-to-noise per coherent chunk is $18/\sqrt{512} = 0.8$.

To interpret the results in Fig. 1, we will compare to analytic estimates of some other search methods.

¹ Our fiducial survey parameters are not completely consistent, since a binary system with $|\ddot{\Phi}| \sim 10^{-4}$ would not be fit by a constant-acceleration model over observation time 3.7 hours. In future work, we plan to extend our methods to searches more complex than a constant-acceleration search. In this paper, the fiducial survey is just intended as a way of comparing our method to its alternatives in a computation-limited regime.

First, we estimate the threshold SNR for a full coherent search as follows. We assume that in the absence of any pulsar signal, $\hat{\mathcal{E}}$ is an independent unit Gaussian for each of the $S \approx 2 \times 10^{13}$ models in the search space. For the coherent search to succeed, the SNR of the pulsar should be comparable to the expected maximum of all S Gaussians. This criterion gives the following estimate:

$$\frac{1}{S} \approx \frac{1}{(2\pi)^{1/2}} \int_r^\infty e^{-x^2/2} dx = \frac{1}{2} \operatorname{erfc} \left(\frac{r}{\sqrt{2}} \right) \quad (26)$$

which gives $r = 7.4$ for $S = 2 \times 10^{13}$. Remarkably, from Fig. 1 we see that the 50% detection threshold for hierarchical search with $N_c = 64$ is $r \approx 9$, which is not much worse! In this case, the hierarchical search does not lose much optimality, but should be faster than a full coherent search by a factor $\approx 64^2 = 4096$.

The analysis in the previous paragraph does not account for the timestream being one of many trial (beam, DM) pairs in a full survey. In a survey, the detection threshold must be set higher than 7.4σ , in order to avoid being flooded with false positives from the large number of trials. To obtain an estimate for a realistic full-survey detection threshold, we apply Eq. (26) with S multiplied by an additional factor $(N_{\text{sky}} N_{\text{DM}}) \sim 10^9$. This gives signal-to-noise threshold $r = 9.8\sigma$. It follows from Fig. 1 that a hierarchical search with $N_c = 64$ is nearly equivalent to a full coherent search, since a pulsar which is bright enough to pass the full-survey SNR threshold will be detected by the hierarchical search 85% of the time. A crucial point here is that the last stage of the hierarchical search is a coherent search, so that if the hierarchical search succeeds then its reported SNR is the true coherent SNR $\hat{\mathcal{E}}$ of the pulsar, which can then be compared to the detection threshold of a full coherent search to reject false positives.

Finally, we would like to compare the hierarchical search to an alternative chunk-based method, namely incoherent power spectrum stacking. In this method, we first divide the timestream into chunks whose size is determined by the criterion that the first D^{-1} harmonics of the pulse period drift by < 1 Fourier bin during one chunk. This gives $T_c \approx (2\pi\alpha_{\text{max}}^{-1} D)^{1/2} \approx 80$ sec, or $N_c = 167$ chunks. We next compute the folded power spectrum $P_{\text{folded}}(\omega)$ in each chunk as described in the introduction. Finally, we loop over trial frequency and acceleration parameters, and sum P_{folded} values over chunks to get a search statistic P_{summed} . (We need to loop over trial accelerations since the pulse frequency can drift by more than one Fourier bin between the first and last chunk.)

To estimate the signal-to-noise threshold for incoherent power spectrum stacking, we will model each per-chunk $P_{\text{folded}}(\omega)$ value as a χ^2 random variable with $D^{-1} = 10$ degrees of freedom. A pulsar with total signal-to-noise r will contribute r^2/N_c to P_{folded} . When we sum P_{folded} values over chunks to get P_{summed} , we get a χ^2 random variable with $d = N_c D^{-1} = 1670$ degrees of freedom, and a pulsar contributes r^2 to P_{summed} . We estimate that the number of trial frequencies needed for the incoherent search is $N_\omega \sim \omega_{\text{max}} T_c / (2\pi)$, and the number of trial accelerations is $N_\alpha \sim N_c$. This gives a total number of trials $N_{\text{trials}} = N_\alpha N_\omega \sim 10^7$. To estimate the threshold signal-to-noise r , we require that the expectation value $(d^2 + r)$ of the P_{summed} statistic which contains the pulsar be greater than the expected maximum of N_{trials} χ^2 random variables with d degrees of freedom. This gives the criterion:

$$\frac{1}{N_{\text{trials}}} = \int_{d+r^2}^\infty dx \frac{x^{d/2-1} e^{-x/2}}{2^{d/2} \Gamma(d/2)} \quad (27)$$

where the integrand on the RHS is the PDF of a χ^2 random variable with d degrees of freedom. This criterion gives signal-to-noise threshold $r \approx 18.6$ for incoherent power spectrum stacking. Comparing with Fig. 1, we see that incoherent power spectrum stacking performs substantially worse than the hierarchical search for the same value of N_c . An additional benefit of the hierarchical search is that N_c is a free parameter which can be adjusted to trade off optimality for computational cost.

IX. DISCUSSION

In this paper, we have proposed four new algorithms for pulsar search: an optimal constant-period search (§III), a recursive tree algorithm for coherent constant-acceleration search (§IV), a semicoherent search which combines information from coherent subsearches while preserving as much phase information as possible (§VI), and a hierarchical search which uses a sequence of semicoherent searches to converge to a coherent search (§VII).

The main result is Fig. 1, where we have simulated the hierarchical search as a function of the number of coherent chunks N_c . The computational cost scales roughly as N_c^{-2} , so the free parameter N_c may be chosen based on total computing time available. We have shown that for surprisingly large N_c , the semicoherent search is nearly equivalent

to a full coherent search. Therefore, as an optimization, it may make sense to use $N_c = 32$ or 64 by default, speeding up the search by a factor 1024 or 4096 . If the search is still too expensive, then larger N_c may be used. The search will then be suboptimal, but less so than incoherent power spectrum stacking.

The semicoherent search has a formal interpretation as the optimal search over the fuzzy search space consisting of all phase models whose acceleration $\ddot{\Phi}$ is allowed to wander over a narrow range $(\Delta\alpha)$. Intuitively, the fuzziness of the search space means that the phase model decorrelates on time scales larger than some coherence time $T_c \propto (\Delta\alpha)^{-1/2}$. The semicoherent search is fully coherent on timescales $\lesssim T_c$. On longer time scales, it starts to become suboptimal, but has the important property that the computational cost grows slowly as $\mathcal{O}(T)$ for $T \gtrsim T_c$. One can think of the semicoherent search as a way of “fuzzing out” the constant-acceleration search on long time scales in order to speed it up.

An interesting property of the semicoherent search is that it should detect pulsars which are approximated by, but not strictly modelled by, a constant-acceleration phase model. For example, a binary pulsar may have a complex phase model with many parameters, but as long as its acceleration varies by less than $(\Delta\alpha)$ over the observation, then the semicoherent search should have the same sensitivity to the binary pulsar as it has to a constant-acceleration pulsar. In a hierarchical search, the binary pulsar would show up as a statistically significant peak in the early stages of the hierarchy, which drops out and loses statistical significance in later stages, when the coherence time T_c becomes large enough that the constant-acceleration model is no longer a good fit. In real data, it will be interesting to flag statistically significant hierarchical search dropouts for human inspection, as a general way of finding pulsars which are approximately but not exactly fit by the assumed search model.

Another useful property of the hierarchical search is that if it succeeds in finding the pulsar, then it finds the full phase model and reports the true phase-coherent signal-to-noise $\hat{\mathcal{E}}$. This is important for rejecting false positives, since the $\hat{\mathcal{E}}$ -values for the best candidates which are found by the semicoherent search may be compared to the detection threshold for a full coherent search, even though the coherent search is too expensive to run over the whole parameter space. This property may also make the hierarchical search more RFI-robust than incoherent power spectrum stacking, since weak unmasked RFI events may contribute statistically to the power spectrum, but are unlikely to have the correct phasing to mimic a coherent pulsar signal whose likelihood is narrowly peaked as a function of phase model parameters, beam, and DM.

In the future, we plan to extend this work in several ways.

Our current implementation is not very well optimized, makes simplifying assumptions such as fixed duty cycle, and is missing features such as dedispersion and RFI removal. Because we do not have an optimized implementation, in this paper we have only been able to make rough estimates of computational cost, rather than giving hard timings. We plan to release a public “production” version soon which will address these shortcomings and be applicable to real data.

In this paper, we have only studied the simple case of a constant-acceleration search space. We speculate that the methods in this paper may be more powerful for more complex search spaces. To take a concrete example, consider the case of a low-order polynomial search, say degree 3 or 4. The cost of a full coherent search grows as $\mathcal{O}(T^6)$ or $\mathcal{O}(T^{10})$ and would quickly become prohibitive. However, one should still be able to define a semicoherent search whose cost is the same as a coherent search up to some coherence time $T_c \ll T$, and grows linearly afterwards. The speedup over a coherent search would be even more dramatic than in the constant-acceleration case. One could even imagine a hierarchical search in which the degree of the polynomial increases during the hierarchy, and in late stages of the hierarchy the polynomial fit is replaced by a more detailed model, such as a many-parameter binary system.

An exciting near-term development in radio astronomy will be the advent of large close-packed transit interferometers such as CHIME [9] and HIRAX [10]. These instruments will have mapping speeds hundreds of times larger than existing telescopes, due to a combination of reasonably large total collecting area, and very large numbers of formed beams. However, with existing algorithms it is difficult to take full advantage of this enormous mapping speed to search for pulsars, since the observing time of each sky location is split into noncontiguous daily observations, and computational cost is also a major issue. It will be very interesting to see whether the methods in this paper can be applied to transit telescopes.

As this paper was nearing completion, we became aware of related work in the context of gravitational wave interferometers such as LIGO. A variety of statistics have been proposed which search long timestreams for quasiperiodic signals by combining information from short coherent searches (e.g. [11–15] and references therein). The method we have proposed is conceptually similar but the details are very different. In the future, we plan to compare these algorithms in more detail.

Acknowledgements

We thank Scott Ransom and Ue-Li Pen for discussions. Research at Perimeter Institute is supported by the Government of Canada through Industry Canada and by the Province of Ontario through the Ministry of Research & Innovation. Some computations were performed on the GPC cluster at the SciNet HPC Consortium. SciNet is funded by the Canada Foundation for Innovation under the auspices of Compute Canada, the Government of Ontario, and the University of Toronto. KMS was supported by an NSERC Discovery Grant and an Ontario Early Researcher Award.

-
- [1] J. H. Taylor and J. M. Weisberg, *Astrophys. J.* **253**, 908 (1982).
 - [2] A. Wolszczan and D. A. Frail, *Nature* **355**, 145 (1992).
 - [3] A. G. Lyne *et al.*, *Science* **303**, 1153 (2004), astro-ph/0401086.
 - [4] I. Yusifov and I. Kucuk, *Astron. Astrophys.* **422**, 545 (2004), astro-ph/0405559.
 - [5] S. L. Detweiler, *Astrophys. J.* **234**, 1100 (1979).
 - [6] K. Liu, R. P. Eatough, N. Wex, and M. Kramer, *Mon. Not. Roy. Astron. Soc.* **445**, 3115 (2014), 1409.3882.
 - [7] E. F. Keane *et al.*, *PoS AASKA14*, 040 (2015), 1501.00056.
 - [8] S. M. Ransom, S. S. Eikenberry, and J. Middleditch, *Astron. J.* **124**, 1788 (2002), astro-ph/0204349.
 - [9] K. Bandura *et al.*, *Proc. SPIE Int. Soc. Opt. Eng.* **9145**, 22 (2014), 1406.2288.
 - [10] L. B. Newburgh *et al.*, *Proc. SPIE Int. Soc. Opt. Eng.* **9906**, 99065X (2016), 1607.02059.
 - [11] P. R. Brady and T. Creighton, *Phys. Rev.* **D61**, 082001 (2000), gr-qc/9812014.
 - [12] B. Krishnan *et al.*, *Phys. Rev.* **D70**, 082001 (2004), gr-qc/0407001.
 - [13] C. Cutler, I. Gholami, and B. Krishnan, *Phys. Rev.* **D72**, 042004 (2005), gr-qc/0505082.
 - [14] H. J. Pletsch and B. Allen, *Phys. Rev. Lett.* **103**, 181102 (2009), 0906.0023.
 - [15] V. Dergachev, *Phys. Rev.* **D85**, 062003 (2012), 1110.3297.

Appendix A: “Transposed” coherent search

In this technical appendix, we define the notion of a transposed search. A transposed search $\hat{\mathcal{E}}^\dagger$ is defined for any coherent search algorithm $\hat{\mathcal{E}}$, but not for the semicoherent search $\hat{\mathcal{H}}$. This is because a coherent search statistic $\hat{\mathcal{E}}$ is a linear function of the timestream data t_k , whereas $\hat{\mathcal{H}}$ is nonlinear.

A coherent search algorithm is applied to a length- N timestream t_k , and produces a 3D array $(\hat{\mathcal{E}}t)_{\alpha\omega\phi}$ on a grid of trial accelerations, frequencies, and phases. Since $\hat{\mathcal{E}}$ is linear, we can view it as a linear operator from an N -component vector space to an $(N_\alpha N_\omega N_\phi)$ -component vector space.

Given two timestreams t_k and t'_k , where $k = 1, \dots, N$, we define their dot product by

$$t \cdot t' = \eta^{-2} t_s \sum_{k=1}^N t_k t'_k \quad (\text{A1})$$

This agrees with our previous definition in Eq. (19). Similarly, given two shape- $(N_\alpha, N_\omega, N_\phi)$ arrays $X_{\alpha\omega\phi}$, $X'_{\alpha\omega\phi}$, we define their dot product by:

$$X \cdot X' = \sum_{\alpha=1}^{N_\alpha} \sum_{\omega=1}^{N_\omega} \sum_{\phi=1}^{N_\phi} X_{\alpha\omega\phi} X'_{\alpha\omega\phi} \quad (\text{A2})$$

Given these dot product definitions, the transposed search $\hat{\mathcal{E}}^\dagger$ is defined to be the formal transpose of the operator $\hat{\mathcal{E}}$ in the usual linear algebra sense. More precisely, given a 3D array $X_{\alpha\omega\phi}$, $(\hat{\mathcal{E}}^\dagger X)$ is a length- N timestream satisfying:

$$(\hat{\mathcal{E}}^\dagger X) \cdot t = X \cdot (\hat{\mathcal{E}}t) \quad (\text{A3})$$

for all timestreams t_k . This equation uniquely determines the operator $\hat{\mathcal{E}}^\dagger$ and can be taken as its definition.

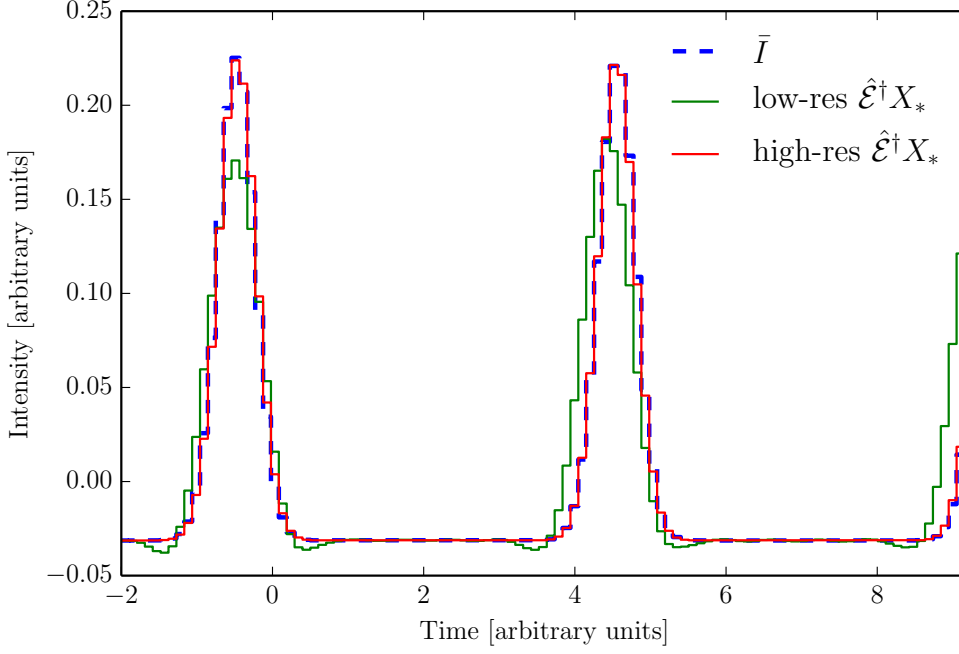


FIG. 2: Visual comparison between the transposed search result $\hat{\mathcal{E}}^\dagger X_*$ for a two-level tree search, and a direct simulation of the pulsar timestream \bar{I} in the time domain, for arbitrarily chosen trial parameters $(\alpha_*, \omega_*, \Phi_*)$. The jaggedness of the curves is due to the discrete timestream sampling. The two $\hat{\mathcal{E}}^\dagger$ curves represent choices of resolution-like parameters, such as spacings in interpolation tables (a complete list of parameters is given in the text). As the resolution increases, $\hat{\mathcal{E}}^\dagger$ converges to \bar{I} precisely, including discrete-sampling artifacts. As explained in the text, this gives a complete test that the tree search is implementing the optimal estimator, and gives a criterion for choosing resolution parameters.

The fast coherent search algorithms from §III, §IV are algorithms which compute $(\hat{\mathcal{E}}t)$ for a fixed timestream t . In this appendix we will show that these algorithms can be formally transposed, to give algorithms for computing $(\hat{\mathcal{E}}^\dagger X)$ for a fixed input array X . First we explain why the transposed search $\hat{\mathcal{E}}^\dagger$ is useful.

Suppose we apply the transpose search operator $\hat{\mathcal{E}}^\dagger$ to a “singleton” array X_* , that is a 3D array whose entries are all zero, except for a single entry which is equal to 1. Let $(\alpha_*, \omega_*, \Phi_*)$ be the trial parameters corresponding to the nonzero entry. Now consider Eq. (A3) defining $\hat{\mathcal{E}}^\dagger$, specialized to the case where $X = X_*$ is a singleton array. The quantity $(X_* \cdot (\hat{\mathcal{E}}t))$ appearing on the RHS is simply the coherent statistic $\hat{\mathcal{E}}t$ evaluated at trial parameters $(\alpha_*, \omega_*, \Phi_*)$. Thus we can write:

$$(\hat{\mathcal{E}}^\dagger X_*) \cdot t = (\hat{\mathcal{E}}t)_{\alpha_*, \omega_*, \Phi_*} \quad (\text{A4})$$

On the other hand, from the definition of $\hat{\mathcal{E}}$ in Eq. (5), the RHS is also equal to $(t \cdot \bar{I}_{\alpha_*, \omega_*, \Phi_*})$, the dot product of t with the normalized signal timestream $\bar{I}_{\alpha_*, \omega_*, \Phi_*}$ corresponding to pulsar parameters $(\alpha_*, \omega_*, \Phi_*)$. Therefore, $(\hat{\mathcal{E}}^\dagger X_*) \cdot t = \bar{I}_{\alpha_*, \omega_*, \Phi_*} \cdot t$. Since this applies to any timestream t , we must have

$$\hat{\mathcal{E}}^\dagger X_* = \bar{I}_{\alpha_*, \omega_*, \Phi_*} \quad (\text{A5})$$

We have now derived a key property of $\hat{\mathcal{E}}^\dagger$. When $\hat{\mathcal{E}}^\dagger$ is applied to a singleton array X_* , the quantity $(\hat{\mathcal{E}}^\dagger X_*)$ is the normalized signal timestream $\bar{I}_{\alpha_*, \omega_*, \Phi_*}$ associated pulsar parameters $(\alpha_*, \omega_*, \Phi_*)$ of the singleton array.

We have found this property to be of great practical use when testing our fast coherent search code. First, we check that our implementations of $t \rightarrow \hat{\mathcal{E}}t$ and $X \rightarrow \hat{\mathcal{E}}^\dagger X$ are consistent, by checking that Eq. (A3) is satisfied for randomly generated pairs (X, t) . Then we choose pulsar parameters $(\alpha_*, \omega_*, \Phi_*)$ and verify that $\hat{\mathcal{E}}^\dagger X_*$ agrees with a direct simulation of the pulsar in the time domain. This comparison is shown visually in Fig. 2. Taken together, these tests give a complete test that the coherent search $t \rightarrow \hat{\mathcal{E}}t$ is correctly implemented and optimal. We use this to test the constant-period search algorithm from §III and the constant-acceleration tree search from §IV.

Note that $\hat{\mathcal{E}}^\dagger X_*$ only agrees with the simulated pulsar timestream \bar{I} in the limit where all resolution-like parameters, for example the spacings of trial parameter grids throughout the tree recursion, are chosen to be high. Away from the high-resolution limit, there is a difference: \bar{I} is the true pulsar timestream which appears in the optimal estimator $\hat{\mathcal{E}}_{\text{opt}} = (t \cdot \bar{I})$, and $\hat{\mathcal{E}}^\dagger X_*$ is the effective pulsar timestream which replaces it in the estimator $\hat{\mathcal{E}}$ due to finite resolution artifacts. The difference between \bar{I} and $\hat{\mathcal{E}}^\dagger X_*$ quantifies the suboptimality due to finite resolution. Therefore, another use for the transposed search is that it gives us a formal criterion for choosing resolution-like parameters. We simply increase resolution until $\hat{\mathcal{E}}^\dagger X_*$ converges to \bar{I} , for a few random choices of X_* . Here is a complete list of resolution-like parameters, and default settings that we obtained using this procedure:

- The spacing $\Delta\alpha$ between trial α values. We find that $\Delta\alpha = 100(D/T_c^2)$ is 3% suboptimal, and $\Delta\alpha = 70(D/T_c^2)$ is 0.5% suboptimal. Here, T_c is the timestream chunk size at the current tree resolution.
- The spacing $\Delta\omega$ between trial ω values. We find that $\Delta\omega = 14(D/T_c)$ is 4% suboptimal, and $\Delta\omega = 10(D/T_c)$ is 0.7% suboptimal.
- The number of trial ϕ values N_Φ . We find that $N_\Phi = (3/2)D^{-1}$ is 3% suboptimal, and $N_\Phi = 2D^{-1}$ is 0.7% suboptimal.
- The threshold timestream size T_0 for switching to a constant-period search, ending the recursion in the tree algorithm. We find that $T_0 = 5(D/|\alpha|_{\text{max}})^{1/2}$ is 4% suboptimal, and $T_0 = 3(D/|\alpha|_{\text{max}})^{1/2}$ is 0.6% suboptimal. Here, $|\alpha|_{\text{max}}$ is the maximum value of $|\alpha|$ in the search.
- The amount of zero-padding used in the bottom-level timestream FFT (Eq. (15)). We find that zero-padding by a factor of two is only 0.1% suboptimal.

In all cases except the last, we have given one value which is a few percent suboptimal, and a more conservative value which is $< 1\%$ suboptimal. We have used the more conservative values as defaults throughout the paper. When we set all parameters to their defaults, we find (by comparing $\hat{\mathcal{E}}^\dagger X_*$ to \bar{I}) that the total suboptimality is 6%. This includes the cumulative effect of multiple interpolations in the tree algorithm.

To conclude this appendix, we explain our algorithm for computing $(\hat{\mathcal{E}}^\dagger X)$ from a 3D array $X_{\alpha\omega\Phi}$.

Let us start with the simplest case, namely the constant-period search from §III. Recall our algorithm for computing $\hat{\mathcal{E}}$. First, we zero-pad the timestream t_k and take its FFT, to obtain the Fourier transform \tilde{t}_j at frequencies $\omega = j\omega_f$, where $\omega_f = 2\pi/T_{\text{padded}}$ is the fundamental frequency of the padded timestream. Let us write this step as:

$$\tilde{t}_j = t_s e^{ij\omega_f(t_s/2 - T/2)} \sum_k e^{ij\omega_f k t_s} t_k \quad (\text{A6})$$

where we have included the prefactor $t_s e^{ij\omega_f(t_s/2 - T/2)}$ for consistency with the previous definition in Eq. (15). Second, we interpolate \tilde{t} to arbitrary frequencies ω . We write this step as:

$$\tilde{t}(\omega) = \sum_j W_j(\omega) \tilde{t}_j \quad (\text{A7})$$

where $W_j(\omega)$ is a sparse interpolation kernel (we have used cubic interpolation throughout this paper). Third, we compute $\hat{\mathcal{E}}$ using Eq. (16), which we repeat here in index notation:

$$(\hat{\mathcal{E}}t)_{\omega\phi} = \frac{1}{A(\omega)^{1/2}\eta^2} \sum_n \rho_n j_0\left(\frac{n\omega t_s}{2}\right) e^{in\phi} \tilde{t}(n\omega) \quad (\text{A8})$$

Given the chain of steps (A6)–(A8) defining $\hat{\mathcal{E}}$, how do we compute $\hat{\mathcal{E}}^\dagger$? We take the defining equation $(\hat{\mathcal{E}}^\dagger X) \cdot t = X \cdot (\hat{\mathcal{E}}t)$ and plug in the above equations for $\hat{\mathcal{E}}t$.

$$\begin{aligned} (\hat{\mathcal{E}}^\dagger X) \cdot t &= X \cdot (\hat{\mathcal{E}}t) \\ &= \sum_{\omega\phi} X_{\omega\phi} \frac{1}{A(\omega)^{1/2}\eta^2} \sum_n \rho_n j_0\left(\frac{n\omega t_s}{2}\right) e^{in\phi} \sum_j W_j(n\omega) t_s e^{ij\omega_f(t_s/2 - T/2)} \sum_k e^{ij\omega_f k t_s} t_k \\ &= \eta^{-2} t_s \sum_k t_k \left[\sum_j e^{ij\omega_f k t_s} e^{ij\omega_f(t_s/2 - T/2)} \sum_{\omega n} W_j(n\omega) \frac{1}{A(\omega)^{1/2}} j_0\left(\frac{n\omega t_s}{2}\right) \sum_\phi e^{in\phi} X_{\omega\phi} \right] \end{aligned} \quad (\text{A9})$$

where in the second line we have plugged in Eqs. (A6)–(A8), and in the third line we have rearranged. In this form, we can read off a formula for $(\hat{\mathcal{E}}^\dagger X)_k$: it is the expression in brackets in the last line. Therefore, $(\hat{\mathcal{E}}^\dagger X)_k$ can be computed from $X_{\omega\phi}$ by the following steps.

$$\begin{aligned}\tilde{X}_{\omega n} &= \frac{1}{A(\omega)^{1/2}} \rho_n j_0 \left(\frac{n\omega t_s}{2} \right) \sum_{\phi} X_{\omega\phi} e^{in\phi} \\ \tilde{E}_j &= \sum_{\omega n} W_j(n\omega) X_{\omega n} \\ (\hat{\mathcal{E}}^\dagger X)_k &= \sum_j e^{ij\omega_f k t_s} e^{ij\omega_f (t_s/2 - T/2)} \tilde{E}_j\end{aligned}\tag{A10}$$

This sequence of steps gives a fast algorithm for computing $(\hat{\mathcal{E}}^\dagger X)$, with the same computational cost as our fast algorithm for computing $(\hat{\mathcal{E}}t)$.

In fact, one can see that the steps (A10)–(A10) in the algorithm for computing $\hat{\mathcal{E}}^\dagger$ are simply the formal transposes of the steps (A6)–(A8), in the reverse order. This makes sense, because if a linear operator factors as $\hat{\mathcal{E}} = \hat{\mathcal{E}}_1 \hat{\mathcal{E}}_2 \hat{\mathcal{E}}_3$, then its transpose also factors as $\hat{\mathcal{E}}^\dagger = \hat{\mathcal{E}}_3^\dagger \hat{\mathcal{E}}_2^\dagger \hat{\mathcal{E}}_1^\dagger$. Given code for computing $t \rightarrow \hat{\mathcal{E}}t$, one can write code for computing $X \rightarrow (\hat{\mathcal{E}}^\dagger X)$ by a fairly mechanical process, by reversing the sequence of steps and replacing each step by its formal transpose. Note that the formal transpose of an FFT is another FFT (e.g. Eqs. (A6), (A10) are transposes). The formal transpose of an interpolation operation, which converts regularly spaced samples of a function to irregularly spaced samples, is a gridding operation which converts irregularly spaced samples to regularly spaced ones (e.g. Eqs. (A7), (A10) are transposes).

This concludes our algorithm for computing $X \rightarrow \hat{\mathcal{E}}^\dagger X$, in the case where $\hat{\mathcal{E}}$ is the fast constant-period search from §III. We now consider the case of the constant-acceleration tree search from §IV. We outline the steps since the ideas are similar to the constant-period case.

If the depth of the recursive tree search is d , then $\hat{\mathcal{E}}$ factors as $\hat{\mathcal{E}} = \hat{\mathcal{E}}_d \cdots \hat{\mathcal{E}}_1 \hat{\mathcal{E}}_0$, where the operator $\hat{\mathcal{E}}_0$ is 2^d copies of the constant-period search, and $\hat{\mathcal{E}}_i$ is 2^{d-i} copies of a tree merge operator $\hat{\mathcal{T}}$, which merges two arrays of shape $(N_\alpha^{\text{in}}, N_\omega^{\text{in}}, N_\phi)$ to a single array of shape $(N_\alpha^{\text{out}}, N_\omega^{\text{out}}, N_\phi)$. Since $\hat{\mathcal{E}}^\dagger = \hat{\mathcal{E}}_0^\dagger \hat{\mathcal{E}}_1^\dagger \cdots \hat{\mathcal{E}}_d^\dagger$, and we have already shown how to compute $\hat{\mathcal{E}}_0^\dagger$, it suffices to compute $\hat{\mathcal{T}}^\dagger$.

Since $\hat{\mathcal{T}}$ is an interpolation operator, its transpose $\hat{\mathcal{T}}^\dagger$ is just the gridding operator with the same weights. To write this out in more detail, let us represent $\hat{\mathcal{T}}$ as an operator which operates on an array X_{abcs} , where $1 \leq \alpha \leq N_\alpha^{\text{in}}$, $1 \leq \omega \leq N_\omega^{\text{in}}$, $1 \leq \phi \leq N_\phi^{\text{in}}$, and the index s is ± 1 . Its output is an array $(\hat{\mathcal{T}}X)_{\alpha\omega\phi}$, where $1 \leq \alpha \leq N_\alpha^{\text{out}}$, $1 \leq \omega \leq N_\omega^{\text{out}}$, and $1 \leq \phi \leq N_\phi^{\text{out}}$. Note that we represent “input” indices with roman letters, and “output” indices with greek letters. From Eq. (18), the tree interpolation $\hat{\mathcal{T}}$ can be written:

$$(\hat{\mathcal{T}}X)_{\alpha\omega\phi} = \frac{1}{\sqrt{2}} \sum_{abcs} W_a(\alpha) W_b \left(\omega + \frac{s\alpha T}{4} \right) W_c \left(\phi + \frac{s\omega T}{4} \right) X_{abcs}\tag{A11}$$

where W_a, W_b, W_c are interpolation weights. Therefore, the transpose gridding operation $\hat{\mathcal{T}}^\dagger$ is:

$$(\hat{\mathcal{T}}^\dagger X)_{abcs} = \frac{1}{\sqrt{2}} \sum_{\alpha\omega\phi} W_a(\alpha) W_b \left(\omega + \frac{s\alpha T}{4} \right) W_c \left(\phi + \frac{s\omega T}{4} \right) X_{\alpha\beta\gamma}\tag{A12}$$