

Bonsai: A GPU Tree-Code

J. Bédorf,¹ E. Gaburov,^{1,2} and S. Portegies-Zwart¹

¹*Leiden Observatory, Leiden University, P.O. Box 9513, 2300 RA Leiden, The Netherlands*

²*Northwestern University, 2131 Tech Drive, Evanston 60208, IL, USA*

Abstract. We present a gravitational hierarchical N -body code that is designed to run efficiently on Graphics Processing Units (GPUs). All parts of the algorithm are executed on the GPU which eliminates the need for data transfer between the Central Processing Unit (CPU) and the GPU. Our tests indicate that the gravitational tree-code outperforms tuned CPU code for all parts of the algorithm and show an overall performance improvement of more than a factor 20, resulting in a processing rate of more than 2.8 million particles per second.

1. Introduction

A popular method for simulating gravitational N -body systems is the hierarchical tree-code algorithm originally introduced by Barnes & Hut (1986). This method reduces the computational complexity of the simulation from $O(N^2)$ to $O(N \log N)$ per crossing time. The former, though computationally expensive, can easily be implemented in parallel for many particles. The latter requires more communication and book keeping when developing a parallel method. Still for large number of particles ($N \gtrsim 10^5$) hierarchical* methods are more efficient than brute force methods. Currently parallel octree implementations are found in a wide range of problems, such as: self gravitating systems, smoothed particle hydrodynamics, molecular dynamics, clump finding, ray tracing and voxel rendering. All of these problems require a high amount of computation time. For high resolution simulations ($N \gtrsim 10^5$) 1 Central Processing Unit (CPU) is not sufficient, one has to use computer clusters or even supercomputers, both of which are expensive and scarce. A GPU provides an attractive alternative to such systems.

The GPU is a massively parallel processor which is specialised in performing independent parallel computations. For an overview and more details see Portegies Zwart et al. (2007); Hamada & Iitaka (2007); Belleman et al. (2008); Gaburov et al. (2009). In this work we have implemented all the separate parts of the Barnes-Hut tree-code algorithm on the GPU. This includes the tree-construction and computation of the tree-properties (multipole moments). By doing so we remove the need to communicate large amounts of data between the CPU and GPU. Since there is no time lost by CPU-GPU

* Tree data-structures are commonly referred to as hierarchical data-structures. In this work we use an octree data-structure.

communication we can make optimal use of the GPU in a block time-step algorithm. In previous work Gaburov et al. (2010); Hamada & Nitadori (2010) parts of the algorithm were executed on the CPU which only allows for efficient (parallel) execution if shared time-steps are used.

Full implementation details, performance and accuracy characteristics can be found in Bédorf et al. (2011)

2. Results

The algorithms are implemented as part of the gravitational N -body code *Bonsai*[†]. We compare the performance of the GPU algorithms with (optimized) CPU implementations of comparable algorithms. The performance of the tree-traverse depends critically on the multipole acceptance criteria (MAC) (θ) which sets the trade-off between speed and accuracy. In our implementation we use a combination of the method introduced by Barnes (1994) and the method used for tree-traversal on vector machines see Barnes (1994, 1990). To compare the CPU and GPU implementations we measure the wall-clock time for the most time critical parts of the algorithms. For the tree-construction we distinguish three parts; Sorting of the particles along a Space Filling Curve (Morton (1966)) (sorting in the figure), reordering of particle properties based on the sort (moving) and construction of the tree-datastructure (tree-construction). Furthermore, timings are presented for the computation of the multipole moments and tree-traverse. As for hardware we used a Xeon E5620 CPU with 4 physical cores and a NVIDIA GTX480 GPU. The results are presented in Fig. 1.

To measure the scaling of the implemented algorithms we execute simulations using Plummer spheres (Plummer (1915)) with $N = 2^{15}$ (32k) up to $N = 2^{22}$ (4M) particles. We measure the performance of the same algorithms as in the previous paragraph. The results can be found in Fig. 2. The wall-clock time spent in the sorting, moving, tree-construction and multipole computation algorithms scales linearly with N for $N \gtrsim 10^6$. For smaller N , however, the scaling is sub-linear, because the algorithms require more than 10^5 particles to saturate the GPU. In theory the tree-traverse scales as $O(N \log N)$, whereas empirically the wall-clock time scales almost linearly with N . This is explained in the inset of Fig. 2, which shows the average number of interactions per particle during the simulation. The average number of particle-cell interactions doubles between $N \gtrsim 32k$ and $N \lesssim 1M$ and keeps gradually increasing for $N \gtrsim 1M$. This break is not clearly visible in the timing results since for small particle numbers ($N \lesssim 1M$) not all GPU resources are saturated. Finally, more than 90% of the wall-clock time is spent on tree-traversal with $\theta = 0.75$. This allows for block time-step execution where the tree-traverse time is reduced by a factor N/N_{active} , where N_{active} is the number of particles that have to be updated.

3. Discussion and Conclusions

We have presented an efficient gravitational N -body tree-code. In contrast to other existing GPU tree-codes, this implementation is executed completely on the GPU. On

[†]The code is publicly available at:
<http://castle.strw.leidenuniv.nl/software.html>




Figure 1. Wall-clock time spent by the CPU and the GPU on various primitive algorithms. The bars show the time spent on the five selected sections of code. The results indicate that our GPU code outperforms the CPU code on all fronts and is between 2 and 30 times faster. Note that the y-axis is in logscale. (Timings using a 2^{20} million body Plummer sphere with $\theta = 0.75$)




Figure 2. The wall-clock time spent by various parts of the program versus the number of particles N . We used Plummer models as initial conditions and varied the number of particles over two orders of magnitude. The scaling of the tree-walk is between $O(N)$ (shown with the black solid line) and the theoretical $O(N \log N)$ and is due to the average number of interactions staying roughly constant (see inset). The asymptotic complexity of the tree-construction approaches $O(N)$, as expected, since all the constituent primitives share the same complexity. The timings are from the GTX480 GPU with $\theta = 0.75$.

a GTX480 the number of particles processed per unit time is 2.8 million particles per second with $\theta = 0.75$. This allows us to routinely carry out simulations on the GPU. Since the current version can only use 1 GPU, the limitation is the amount of memory. For 5 million particles ~ 1 gigabyte of GPU memory is required.

Even though the sorting, moving and tree-construction parts of the code take up roughly 10% of the execution time in the presented timings, these methods do not have to be executed during each time-step when using the block time-step method. It is sufficient to only recompute the multipole moments of tree-cells that have updated child particles, and only when the tree-traverse shows a considerable decline in performance does the complete tree-structure has to be rebuild. This decline is the result of inefficient memory reads and an increase of the average number of particle-cell and particle-particle interactions. This quantity increases because the tree-cell size increases, which causes more cells to be opened by the MAC.

Although the implemented algorithms are designed for a shared-memory architecture, they can be used to construct and traverse tree-structures on parallel GPU clusters using the methods described in Warren & Salmon (1993); Dubinski (1996). Furthermore, in case of a parallel GPU tree-code, the CPU can exchange particles with the other nodes, while the GPU is traversing the tree-structure of the local data, making it possible to hide most of the communication time.

The implemented algorithms are not limited to the evaluation of gravitational forces, but can be applied to a variety of problems, such as neighbour search, clump finding algorithms, fast multipole method and ray tracing. In particular, it is straightforward to implement Smoothed Particle Hydrodynamics in such a code, therefore having a self-gravitating particle based hydrodynamics code implemented on the GPU.

Acknowledgments. This work is supported by NOVA and NWO grants (#639.073-.803, #643.000.802, and #614.061.608, VICI #643.200.503, VIDI #639.042.607). The authors would like to thank Massimiliano Fatica and Mark Harris of NVIDIA for the help with getting the code to run on the Fermi architecture

References

- Barnes, J. 1994, *Computational Astrophysics* (Berlin: Springer-Verlag)
- Barnes, J., & Hut, P. 1986, *Nat*, 324, 446
- Barnes, J. E. 1990, *Journal of Computational Physics*, 87, 161
- Bédorf, J., Gaburov, E., & Portegies Zwart, S. 2011, *ArXiv e-prints*. 1106.1900
- Belleman, R. G., Bédorf, J., & Portegies Zwart, S. F. 2008, *New Astronomy*, 13, 103
- Dubinski, J. 1996, *New Astronomy*, 1, 133. [arXiv:astro-ph/9603097](#)
- Gaburov, E., Bédorf, J., & Portegies Zwart, S. 2010, in *International Conference on Computational Science 2010* (Elsevier)
- Gaburov, E., Harfst, S., & Portegies Zwart, S. 2009, *New Astronomy*, 14, 630
- Hamada, T., & Iitaka, T. 2007, *ArXiv Astrophysics e-prints*. [arXiv:astro-ph/0703100](#)
- Hamada, T., & Nitadori, K. 2010, in *Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis* (Washington, DC, USA: IEEE Computer Society), SC '10, 1
- Morton, G. M. 1966, *A computer Oriented Geodetic Data Base; and a New Technique in File Sequencing*, Tech. rep., Ottawa, Canada: IBM Ltd.
- Plummer, H. C. 1915, *MNRAS*, 76, 107
- Portegies Zwart, S. F., Belleman, R. G., & Geldof, P. M. 2007, *New Astronomy*, 12, 641
- Warren, M. S., & Salmon, J. K. 1993, in *Supercomputing '93: Proceedings of the 1993 ACM/IEEE conference on Supercomputing* (New York, NY, USA: ACM), 12